

Cache-aware data structures for packet forwarding tables on general purpose CPUs: project description

Maksim Yegorov

2018-02-13 Tue 01:47 PM

context

- ▶ network router receives a packet
- ▶ inspects destination ip
- ▶ looks up by prefix
- ▶ forwards to outgoing interface for next hop




Figure 1: forwarding; from Kurose & Ross, *Computer Networking*

constraints

- ▶ time: key-value lookup
- ▶ space: table encoding
- ▶ hardware cost, power efficiency etc.

estimate throughput demand

ballpark estimate:

- ▶ ≈ 50 Gbps router $\equiv 75$ Mpps
- ▶ varying prefix length, conservatively estimate $\times 32$
- ▶ $\rightarrow 2.5e9$ lookups per second
- ▶ @4GHz single CPU $\rightarrow 50$ CPU cycles per packet
- ▶ ignore other overhead (routing table etc.) as performed by dedicated CPU/memory

FIB table size




Figure 2: IPv4 table; source: <http://bgp.potaroo.net/>




Figure 3: IPv6 table

time and space requirements

space

@ 100B per key-value pair, $1e6$ entries \rightarrow 100MB hash table
(vs. 4MB Bloom filter)

time

- ▶ @10 CPU cycles per byte of hash \rightarrow 40 cycles / 32b hash
- ▶ fetching from memory: expensive

cache loads and misses

Cache type	Access time (cycles)	Cache size (C)	Assoc. (E)	Block size (B)	Sets (S)
L1 i-cache	4	32 KB	8	64 B	64
L1 d-cache	4	32 KB	8	64 B	64
L2 unified cache	10	256 KB	8	64 B	512
L3 unified cache	40–75	8 MB	16	64 B	8,192

Figure 6.39 Characteristics of the Intel Core i7 cache hierarchy.

Miss penalty. Any additional time required because of a miss. The penalty for L1 misses served from L2 is on the order of 10 cycles; from L3, 50 cycles; and from main memory, 200 cycles.

Figure 4: from Bryant & O'Hallaron, *Computer Systems*

... ouch!

existing solutions

- ▶ well researched area
- ▶ traditionally the lookup is performed in highly optimized hardware, using compactly encoded datastructures (trie / Bloom filter and derivatives)
- ▶ TCAM: mask and compare incoming packet in parallel for all prefix lengths
- ▶ SRAM: a couple orders of magnitude cheaper than TCAM

recent proposals

- ▶ state of the art multithreading, algorithms and vector packet processing make forwarding feasible on commodity hardware
- ▶ advantages: cost, added flexibility/programmability, lower barrier to entry (not requiring advanced manufacturing capabilities)
- ▶ **bottleneck**: memory access
- ▶ challenge: how is it even possible?

cache optimization strategies

- ▶ design for sequential reads & writes
- ▶ minimize misses by prefetching in cache strategically
- ▶ helper threads
- ▶ datastructures designed for concurrency
- ▶ arrange elements with memory alignment in mind
- ▶ split tables into long-tail low-fidelity and short-tail high-fidelity:
 - ▶ avoid having the more frequent prefix lengths require additional false-positive test
 - ▶ design split table size to (partially) fit in L3 cache

plan

- ▶ **milestone 1:** review data structures, memory hierarchy & cache optimization literature
- ▶ **milestone 2:** implement filters (Bloom, cuckoo) and hash table (low- vs. high-fidelity versions) with and without cache optimization
- ▶ **milestone 3:** analyze patterns of cache hits and misses, run tests, and benchmark