

Cache-aware data structures for packet
forwarding tables on general purpose CPUs:
Milestone 2

Maksim Yegorov

2018-03-22 Thu 01:25 PM

presentation highlights

1. What is the longest prefix match problem?
2. What do I propose to do?
 - ▶ theme: improve upon linear search
 - ▶ algorithm: optimal binary search tree (Knuth)
 - ▶ algorithm: IP lookup in the table
 - ▶ algorithm: build the lookup table
 - ▶ testing program
3. Goals for milestone 3

context

- ▶ task: match destination IP 1.2.3.4 from an incoming packet
- ▶ Classless Inter Domain Routing (CIDR) tradeoff: use limited IP address space efficiently at the cost of lookup complexity
- ▶ router forwarding table

Network	Next Hop
...	
*> 1.2.0.0/16	203.133.248.254
...	
*> 1.2.3.0/24	202.12.28.1
...	

context (cont.)

Constraints:

- ▶ process packets at line speed or drop packets
- ▶ target programmable routers on conventional CPUs

Solutions:

- ▶ tries
- ▶ hash tables
- ▶ filters (Bloom, cuckoo)

Bloom filter

- ▶ Original network applications of the BF made use of dedicated hardware that enables searching all prefix lengths in parallel
- ▶ Later adaptations in software use linear search




Figure 1: Bloom filter (source: Wikipedia)

```
*> 1.2.0.0/16      x
*> 1.2.3.0/24      y
> 1.2.3.4          w
```

guided search

- ▶ BF involves multiple independent hash functions to insert or look up a prefix
- ▶ Use the first hash function, $hash_1$ to serve as a marker:
 - ▶ hit \rightarrow look right
 - ▶ miss \rightarrow look left
- ▶ for subprefixes in the “happy” search path, check up to $\log(N)$ hashes during the search (in case of a balanced tree, N is the number of valid prefix lengths)
- ▶ will talk more about signaling later

(optimal) binary search tree




Figure 2: weighted tree cost; source: Knuth

(optimal) binary search tree (cont.)

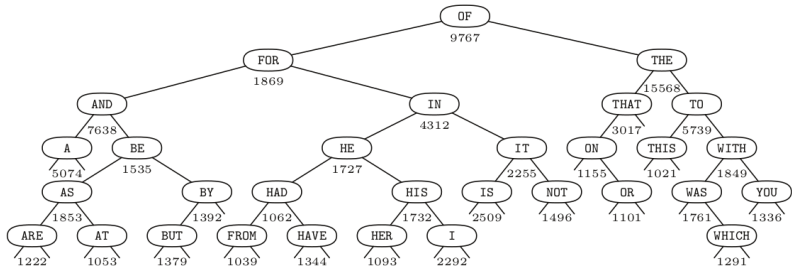


Figure 3: example optimal BST; source: Knuth

use case for BST optimization




Figure 4: prefix length distribution in backbone routers ca. 2000 (source: Waldvogel et al.)

(optimal) binary search tree: algorithm sketch (Knuth)

Optimal substructure: each subtree T' of an optimal BST T must itself be optimal.

When adding a root node to optimal BST subtrees, the total tree cost becomes:

$$e_{i,j} = w_r + (e_{i,r-1} + w_{i,r-1}) + (e_{r+1,j} + w_{r+1,j})$$

Seek the root that minimizes the tree cost:

$$r_{i,j} = \operatorname{argmin}_{i \leq r \leq j} \{e_{i,r-1} + w_{i,j} + e_{r+1,j}\}$$

what's an optimal weighting for prefix distribution?

- ▶ ultimately depends on the traffic (unknown)
- ▶ experiment with:
 - ▶ balanced tree
 - ▶ weighting based on the fraction of IP address space covered by prefixes of given length
 - ▶ weighting based on the fraction of prefixes in table of given length

algorithm: IP lookup in the table




Figure 5: tree traversal

algorithm: IP lookup in the table (cont.)

- ▶ Bloom filter uses $hash_{1..k}$ for set membership ($k \approx 20$)
- ▶ $hash_1$ serves as a marker reserved for guiding the search:
 - ▶ miss \rightarrow look left, else goto most recent $prefix_{hit}$, else $prefix_{default}$
 - ▶ hit \rightarrow remember this prefix length, increment $count_{hit}$ & look right
- ▶ if hit with nil right child or if returning to most recent $prefix_{hit}$, decode $hash_{count_{hit}..count_{hit}+n}$ to calculate *best matching prefix* for $prefix_{hit}$, finish matching with $hash_{n+1..k}$
- ▶ edge cases:
 - ▶ if none of $hash_{count_{hit}..count_{hit}+n}$ succeed \rightarrow default path
 - ▶ if all of $hash_{count_{hit}..k}$ succeed \rightarrow found match
- ▶ n is 5 for IPv4, 6 for IPv6 (bits to encode valid prefix lengths, implicitly assuming $n \geq count_{hit}$)

algorithm: IP lookup in the table (cont.)

- ▶ guided search fails on false positives
- ▶ false positive probability is exceedingly small (put otherwise, probability \gg chance if any single hash matches)
- ▶ on false positive, default to linear search of prefix lengths $\leq prefix_{hit}$




Figure 6: false positive probability vs num hash functions (k)

algorithm: build the lookup table

- ▶ idea: precompute traversal paths at insertion
- ▶ compute optimal BST to direct the search
- ▶ sort prefixes prior to inserting into the Bloom filter, used to find the *best matching prefix* via lookup in already built BF
- ▶ insert prefixes in the sorted order, as follows:
 - ▶ find BMP for given prefix
 - ▶ traverse the BST in search of the prefix length, insert $hash_1$ markers for *branch right* up to and inclusive of the prefix length, increment $count_{hit}$
 - ▶ insert prefix into Bloom filter using $hash_{1..k}$ (can't be reduced to allow for linear search all else failing)
 - ▶ encode pointers (binary encoding for *best matching prefix* length) at each ancestor with marker (i.e. each ancestor where we branched right) on the path to this prefix, using $hash_{count_{hit}..count_{hit}+n}$
 - ▶ edge case: ancestor is itself a BMP \rightarrow previously encoded with $hash_{count_{hit}..k}$
 - ▶ take care to increment $count_{hit}$ going down the tree, decrement when going up the tree

testing program: linear vs. guided search

1. metrics: average number of matches per IP + packets (IPs) per second
2. false positive rate:
 - ▶ increase in false positive rate because of inserted pointers (fraction of Bloom filter filled other things constant)
 - ▶ what fraction of prefixes in the table have a shorter *best matching prefix* (other than default)
3. effects of simulated traffic pattern:
 - ▶ randomly generated traffic over all address space (with default route)
 - ▶ randomly generated from address space covered by prefixes (no default route)
 - ▶ traffic correlated with prefix table (prefix length distribution, fraction of address space covered by each prefix)
 - ▶ effects of optimal binary search tree (balanced vs. weighted by covered address space vs. weighted by fraction of prefix length)
4. performance on IPv4 vs IPv6 prefixes

summary

- ▶ tightened the algorithms since last presentation
- ▶ defined the testing program
- ▶ working on implementation

milestone 3

- ▶ complete implementation
- ▶ analyze results from experiments