

Cache-aware data structures for packet
forwarding tables on general purpose CPUs:
Milestone 1

Maksim Yegorov

2018-02-26 Mon 10:29 PM

context




Figure 1: forwarding; from Kurose & Ross, *Computer Networking*

problem: demand $>$ capacity

- ▶ process packets at line speed or drop packets
- ▶ FIB table exhibits polynomial growth
- ▶ target programmable routers on conventional CPUs where:
 - ▶ DRAM latency almost unchanged for decades
 - ▶ workhorse data structures must fit in limited CPU cache
 - ▶ limited budget of CPU cycles per packet

classes of proposed data structures

- ▶ tries: too much overhead to fit in cache; best solutions patented
- ▶ binary search on prefix ranges (Lampson, Srinivasan, Varghese) and prefix lengths (Waldvogel)
- ▶ filters (Bloom, cuckoo)

filters

Opportunities for improvement:

- ▶ lots of lookups required to find longest prefix match:
 - ▶ 8..32 prefix lengths (IPv4)
 - ▶ 19..128 prefix lengths (IPv6)
- ▶ lookup in off-chip hash table is prohibitively expensive (TBD)

take inspiration from binary search

Try to **nudge filter lookup in the right direction** by quasi-false-positives:

Algorithm 1 wishful thinking heuristic

Input: BF, IP, $8 < i < 32$

Output: j (next prefix length to consider) or i (done)

```
1:  $h \leftarrow \text{hash}(\text{IP}[0:i])$ 
2: if BF.lookup( $h$ ) = 1 then
3:   investigate  $j > i$ 
4:   if BF.match(IP[: $i$ ]) then
5:     return  $i$ 
6:   end if
7: else
8:   investigate  $j < i$ 
9: end if
```

false positive probability vs. number of hashes

- ▶ false positive probability is exceedingly small (put otherwise, probability \gg chance if any single hash matches)
- ▶ empirical fact: uncommon that table includes several prefixes within same subsequence




Figure 2: false positive probability vs num hash functions

more plausible heuristic

Proceed in logarithmic stride:

1. When populating the filter, apply $hash_1$ to subprefixes in search path (at most, $\log(N)$ hashes: e.g. for a $prefix_{32}$ of length 32, apply $hash_1$ to $prefix_{\{16,24,28,30\}}$)
2. apply up to $k - \log(N)$ remaining hash functions to the $prefix$ (e.g. apply $hash_{2..(k-4)}$ to $prefix_{32}$ in this example)
3. start lookup on $prefix_{16}$ (or perhaps most frequent prefix length) and climb if at least non-zero match
4. remember:
 - ▶ longest match to date,
 - ▶ computed hashes

edge cases to consider

- ▶ iff given prefix is in fact a set member, we can always ascertain this in $\log(N)$ steps
- ▶ but what is the effect of false positives? may need to backtrack

summary

Recall what we set out to do:

- ▶ lots of lookups required to find longest prefix match:
 - ▶ 8..32 prefix lengths (IPv4)
 - ▶ 19..128 prefix lengths (IPv6)

summary

- ▶ approach lookup in FIB filters in a way similar to hill climbing and binary search based approaches
- ▶ need to experiment, see the effect on false positives, tweak params as needed, see if can parallelize search
- ▶ on a different note: still need to think about off-chip table lookup